

Required Libraries

```
In [22]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
nltk.download('stopwords')
nltk.download('wordnet')
```

[nltk_data] Downloading package stopwords to C:\Users\s monisi
[nltk_data] prabha\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to C:\Users\s monisi
[nltk_data] prabha\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!

Out[22]: True

Load dataset

```
In [23]: file_path = r'C:\Users\s monisi prabha\Downloads\disaster_tweets_data(DS).csv'
data = pd.read_csv(file_path)
```

Handle missing values

```
In [24]: data = data.dropna()
```

Preprocess tweets

```
In [25]: def preprocess_text(text):
text = re.sub(r'http\S+|www\S+|https\S+', '', text, flags=re.MULTILINE) # Remove
text = re.sub(r'@\w+|\#', '', text) # Remove mentions and hashtags
text = re.sub(r'[\W\s]', '', text) # Remove punctuations
text = text.lower() # Convert to Lowercase
tokens = text.split() # Tokenize the text
tokens = [word for word in tokens if word not in stopwords.words('english')] #
lemmatizer = WordNetLemmatizer()
tokens = [lemmatizer.lemmatize(word) for word in tokens] # Lemmatize words
return ' '.join(tokens)
```

Apply preprocessing to the dataset

```
In [26]: data['tweets'] = data['tweets'].apply(preprocess_text)
```

Convert tweets into vectors using CountVectorizer or TF-IDF

Using TF-IDF

```
In [27]: vectorizer = TfidfVectorizer(max_features=5000)
x = vectorizer.fit_transform(data['tweets']).toarray()
```

Split data into training and testing

```
In [28]: y = data['target']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_stat
```

Multinomial Naive Bayes

```
In [29]: nb_model = MultinomialNB()
nb_model.fit(x_train, y_train)
y_pred_nb = nb_model.predict(x_test)
```

Logistic Regression

```
In [30]: lr_model = LogisticRegression(max_iter=1000)
lr_model.fit(x_train, y_train)
y_pred_lr = lr_model.predict(x_test)
```

K-Nearest Neighbors

```
In [31]: knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(x_train, y_train)
y_pred_knn = knn_model.predict(x_test)
```

Evaluate the models

```
In [32]: def evaluate_model(y_test, y_pred, model_name):
    print(f"Model: {model_name}")
    print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
    print("Classification Report:\n", classification_report(y_test, y_pred))
    print("Accuracy: ", accuracy_score(y_test, y_pred))
    print("\n")
```

Evaluate NB

```
In [33]: evaluate_model(y_test, y_pred_nb, "Multinomial Naive Bayes")
```

Model: Multinomial Naive Bayes

Confusion Matrix:

```
[[772 102]
```

```
[205 444]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.88	0.83	874
1	0.81	0.68	0.74	649
accuracy			0.80	1523
macro avg	0.80	0.78	0.79	1523
weighted avg	0.80	0.80	0.80	1523

Accuracy: 0.7984241628365069

Evaluate LR

```
In [34]: evaluate_model(y_test, y_pred_lr, "Logistic Regression")
```

Model: Logistic Regression

Confusion Matrix:

```
[[776 98]
```

```
[211 438]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.89	0.83	874
1	0.82	0.67	0.74	649
accuracy			0.80	1523
macro avg	0.80	0.78	0.79	1523
weighted avg	0.80	0.80	0.79	1523

Accuracy: 0.7971109652002626

Evaluate KNN

```
In [35]: evaluate_model(y_test, y_pred_knn, "KNN")
```

Model: KNN

Confusion Matrix:

```
[[858 16]
```

```
[462 187]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.65	0.98	0.78	874
1	0.92	0.29	0.44	649
accuracy			0.69	1523
macro avg	0.79	0.63	0.61	1523

weighted avg	0.77	0.69	0.64	1523
--------------	------	------	------	------

Accuracy: 0.6861457649376231

Compare accuracies and select the best model

In [36]:

```
models = {  
    "Multinomial Naive Bayes": accuracy_score(y_test, y_pred_nb),  
    "Logistic Regression": accuracy_score(y_test, y_pred_lr),  
    "KNN": accuracy_score(y_test, y_pred_knn)  
}  
  
best_model = max(models, key=models.get)  
print(f"Best model: {best_model} with accuracy {models[best_model]:.4f}")
```

Best model: Multinomial Naive Bayes with accuracy 0.7984